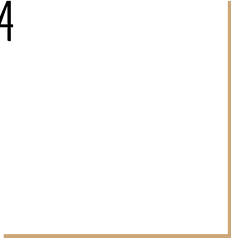




# Multimodal RAG

December 19, 2024

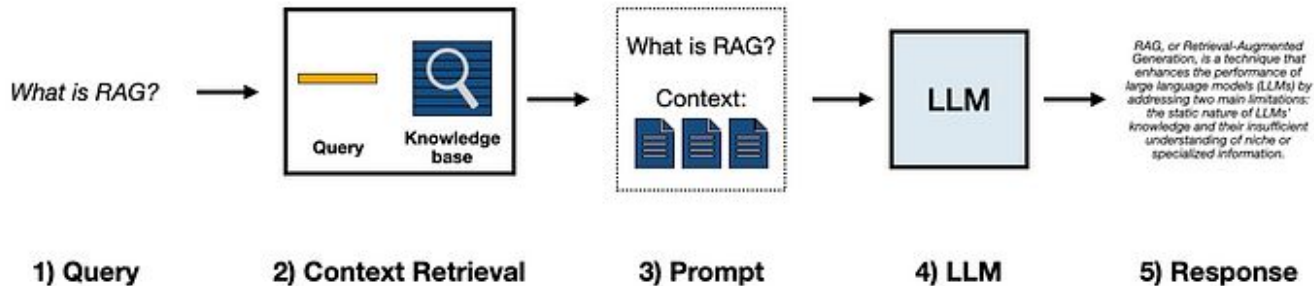


Himanshu Singh Rao  
Math 608  
012629202

# What is Retrieval Augmentation Generation(RAG)?

- LLMs dont have the context of our data, if we give LLMs context from the knowledge base, they become really powerful
- Add your PDFs, add your database, add your knowledge base.
- API call to LLMs = query + my\_context + prompt
- Each time you make a call just pass your context in it

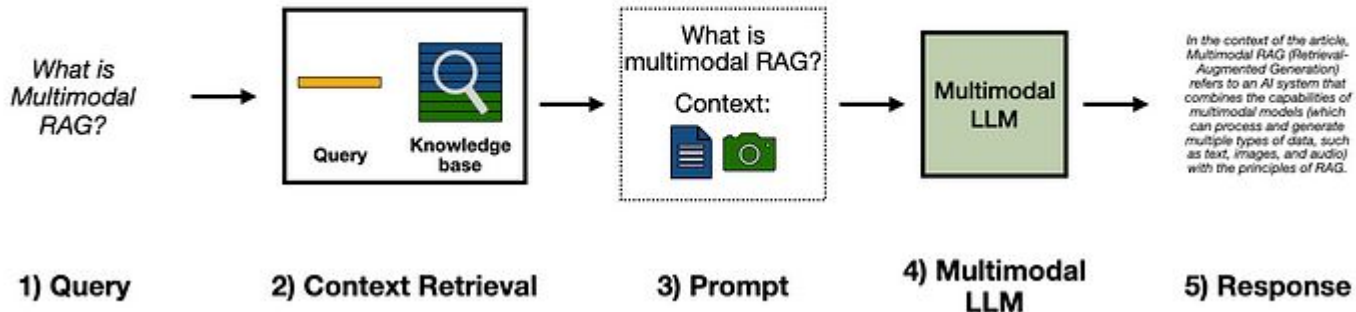
## Basic RAG Design



# What is Multimodal-RAG?

- What if I don't want to make my Llm only text based?
- What if I want my LLM to understand, photo, video, audio and text?
- I want to add images too. How can I do it??
- Multimodal RAG understands all modalities (Image, text, audio, video)

## Basic Multimodal RAG Design



# Why it is Important?

- Chat with Videos
  - Chat with Photos
  - Doing RAG with both Structured/Unstructured data
  - Get full context of pdfs
  - All modalities considered
  - Make really cool AI Applications
-

# 3 Levels of MRAG

---

## Level 1

Translate  
modalities to text.

## Level 2

Text-only retrieval  
+ MLLM

## Level 3

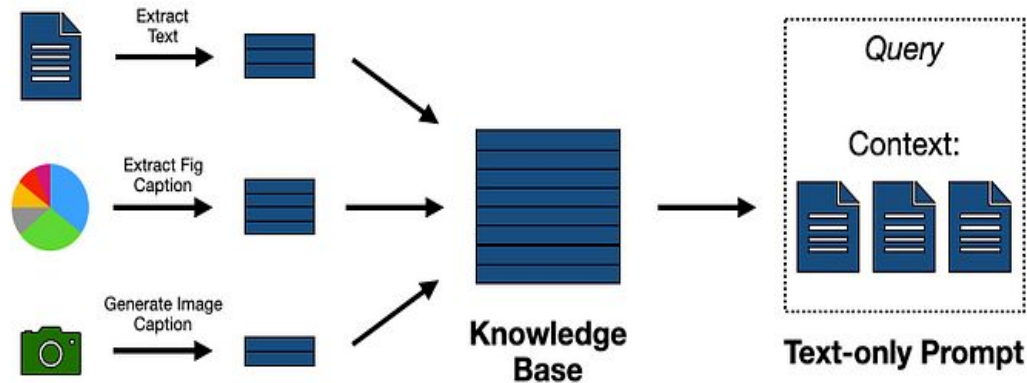
Multimodal  
retrieval + MLLM

# Level 1 - Translate modalities to text

---

- A simple way to make a RAG system multimodal is by translating new modalities to text before storing them in the knowledge base
- This includes tasks like transcribing meeting recordings, generating image captions with MLLMs, or converting tables into readable formats like .csv or .json.).

## Level 1: Translate Everything to Text




# Level 2 - Text-only retrieval + MLLM

---

- Another approach involves generating text representations, like descriptions and meta-tags, for knowledge base items to aid retrieval, while passing the original modality to a multimodal LLM (MLLM). For instance, image metadata is used for retrieval, and the image itself is sent to the model for inference.

## Level 2: Text-only retrieval + MLLM

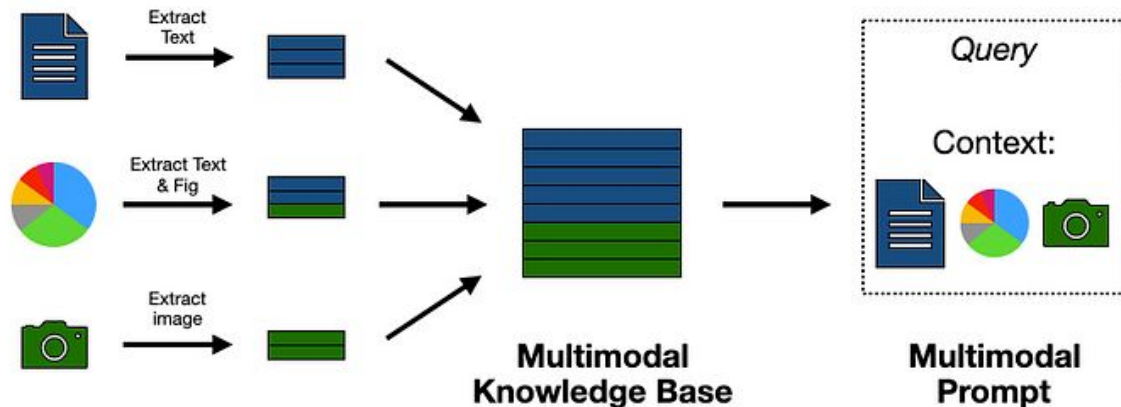


# Level 3 - Multimodal retrieval + MLLM

---

- Use multimodal embeddings to perform multimodal retrieval. This works the same way as text-based vector search, but now the embedding space co-locates similar concepts independent of its original modality. The results of such a retrieval strategy can then be passed directly to a MLLM.

## Level 3: Multimodal retrieval + MLLM





# Implementation/Code walkthrough

---

- Import data
- Multi-vector retriever
- Building a RAG chain
- Querying
- Getting relevant images from the context

The project focuses on searching through a PDF that includes images and tables, with the LLM being able to understand both the content of the images and the text, and tables.

```
import logging
import zipfile
import requests

logging.basicConfig(level=logging.INFO)

data_url = "https://storage.googleapis.com/benchmarks-artifacts/langchain-docs-benchmarking/cj.zip"
result = requests.get(data_url)
filename = "cj.zip"
with open(filename, "wb") as file:
    file.write(result.content)

with zipfile.ZipFile(filename, "r") as zip_ref:
    zip_ref.extractall()

from langchain_community.document_loaders import PyPDFLoader

loader = PyPDFLoader("./cj/cj.pdf")
docs = loader.load()
tables = []
texts = [d.page_content for d in docs]
```

## Importing Data

Loading a pdf with images and tables into the google colab

```
def create_multi_vector_retriever(
    vectorstore, text_summaries, texts, table_summaries, tables, image_summaries, images
):
    """
    Create retriever that indexes summaries, but returns raw images or texts
    """

    # Initialize the storage layer
    store = InMemoryStore()
    id_key = "doc_id"

    # Create the multi-vector retriever
    retriever = MultiVectorRetriever(
        vectorstore=vectorstore,
        docstore=store,
        id_key=id_key,
    )

    # Helper function to add documents to the vectorstore and docstore
    def add_documents(retriever, doc_summaries, doc_contents):
        doc_ids = [str(uuid.uuid4()) for _ in doc_contents]
        summary_docs = [
            Document(page_content=s, metadata={id_key: doc_ids[i]})
            for i, s in enumerate(doc_summaries)
        ]
        retriever.vectorstore.add_documents(summary_docs)
        retriever.docstore.mset(list(zip(doc_ids, doc_contents)))

    # Add texts, tables, and images
    # Check that text_summaries is not empty before adding
    if text_summaries:
        add_documents(retriever, text_summaries, texts)
    # Check that table_summaries is not empty before adding
    if table_summaries:
        add_documents(retriever, table_summaries, tables)
    # Check that image_summaries is not empty before adding
    if image_summaries:
        add_documents(retriever, image_summaries, images)

    return retriever
```

## Creating a retriever

Creating a multivector retriever that stores images, text and table data.

```
# The vectorstore to use to index the summaries
vectorstore = Chroma(
    collection_name="mm_rag_cj_blog",
    embedding_function=VertexAIEmbeddings(model_name="textembedding-gecko@latest"),
)

# Create retriever
retriever_multi_vector_img = create_multi_vector_retriever(
    vectorstore,
    text_summaries,
    texts,
    table_summaries,
    tables,
    image_summaries,
    img_base64_list,
)
```

## Chroma Db Vectorstore

---

Using chroma db Vectorstore to make a multi-vector vector retriever, which is an in-memory store.

```
def multi_modal_rag_chain(retriever):
    """
    Multi-modal RAG chain
    """

    # Multi-modal LLM
    model = ChatVertexAI(
        temperature=0, model_name="gemini-pro-vision", max_output_tokens=1024
    )

    # RAG pipeline
    chain = (
        {
            "context": retriever | RunnableLambda(split_image_text_types),
            "question": RunnablePassthrough(),
        }
        | RunnableLambda(img_prompt_func)
        | model
        | StrOutputParser()
    )

    return chain

# Create RAG chain
chain_multimodal_rag = multi_modal_rag_chain(retriever_multi_vector_img)
```

## Creating RAG Chain

---

Create RAG chain with Google Vertex Api, Langchain and Chroma db.

```
query = "What are the EV / NTM and NTM rev growth for MongoDB, Cloudflare, and Datadog?"  
docs = retriever_multi_vector_img.get_relevant_documents(query, limit=1)
```

```
# We get relevant docs  
len(docs)
```

2

```
plt_img_base64(docs[0])
```

## Querying

We can query the entire PDF, including images, text, and tables.

Company	EV / NTM Rev	EV / 2024 Rev	EV / NTM FCF	NTM Rev Growth	Gross Margin
1 Snowflake	15.5x	13.4x	55x	27%	66%
2 MongoDB	14.6x	12.9x	133x	17%	74%
3 Palantir	14.5x	13.9x	58x	19%	80%
4 Cloudflare	13.4x	12.6x	153x	28%	76%
5 Datadog	13.1x	12.4x	52x	19%	80%
6 CrowdStrike	12.5x	11.1x	37x	31%	74%
7 Adobe	12.3x	11.9x	30x	12%	88%
8 ServiceNow	12.2x	11.6x	38x	21%	79%
9 Samsara	11.8x	10.5x	393x	31%	72%
10 Zscaler	11.8x	10.5x	48x	27%	78%
Average	13.2x	12.1x	100x	23%	77%
Median	12.8x	12.2x	54x	24%	77%
Overall Median	5.0x	4.8x	33.7x	15%	75%
Clouded Judgement		@jaminball			

## Result

The retrieved image based on the query.

```
chain_multimodal_rag.invoke("what is the document about?")
```

'The document is about the valuation of SaaS businesses. It discusses the use of revenue multiples as a short

```
result = chain_multimodal_rag.invoke(query)
print(result)
```

```
| Company | EV/NTM Rev | NTM Rev Growth |
| --- | --- | --- |
| MongoDB | 14.6x | 17% |
| Cloudflare | 13.4x | 28% |
| Datadog | 13.1x | 19% |
```

```
from IPython.display import Markdown as md
md(result)
```

Company	EV/NTM Rev	NTM Rev Growth
MongoDB	14.6x	17%
Cloudflare	13.4x	28%
Datadog	13.1x	19%

---

## Result

We asked LLM a very specific question about the revenue which was present only in the image, and LLM was able to answer it properly.

# Applications that you can build



## Virtual Interviewer

---

An AI virtual interviewer, that you can practice interviews with.



## Calorie Tracker

---

A calorie tracker app that give the calories of the food based on the image provided



## Image Search Engine

---

A search engine that can understand images, and videos.



## Stock Analyzer

---

A stock analyzer that can understands complex tables and charts

# Conclusion

We understood what Multimodal RAG is, its importance, applications, and also examined the implementation of a Multimodal RAG chain.